

PREDICTIVE MODEL

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
import tensorflow as tf
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.model_selection import train_test_split
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import SelectKBest, chi2
from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

```
import os
os.environ['OMP_NUM_THREADS']='1'
```

```
from google.colab import drive
drive.mount('/content/drive/')
```

Mounted at /content/drive/

```
fileName = '/content/drive/MyDrive/LBG Data.csv'
```

```
data = pd.read_csv(fileName)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18324 entries, 0 to 18323
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    18324 non-null  int64
1   addr_state                           18324 non-null  object
2   annual_inc                           18324 non-null  float64
3   emp_length                           17150 non-null  float64
4   emp_title                            17042 non-null  object
5   home_ownership                       18324 non-null  object
6   installment                          18324 non-null  float64
7   loan_amnt                            18324 non-null  int64
8   purpose                              18324 non-null  object
9   term                                 18324 non-null  int64
10  int_rate                             18324 non-null  float64
11  avg_cur_bal                          17758 non-null  float64
12  inq_last_12m                         9395 non-null   float64
13  max_bal_bc                           9395 non-null   float64
14  mo_sin_old_il_acct                   17192 non-null  float64
15  mo_sin_old_rev_tl_op                 17760 non-null  float64
16  mo_sin_rcnt_rev_tl_op                17760 non-null  float64
17  mo_sin_rcnt_tl                       17760 non-null  float64
18  mort_acc                             17926 non-null  float64
19  mths_since_last_delinq               9276 non-null   float64
20  num_bc_tl                            17760 non-null  float64
21  num_il_tl                            17760 non-null  float64
22  num_op_rev_tl                        17760 non-null  float64
23  num_tl_90g_dpd_24m                  17760 non-null  float64
```

```


24  num_tl_op_past_12m      17760 non-null float64
25  open_acc                18324 non-null int64
26  percent_bc_gt_75        17714 non-null float64
27  pub_rec_bankruptcies    18324 non-null int64
28  total_acc               18324 non-null int64
29  total_bal_ex_mort        17926 non-null float64
30  loan_status             18324 non-null object
dtypes: float64(20), int64(6), object(5)
memory usage: 4.3+ MB

```

```
data.describe()
```

	id	annual_inc	emp_length	installment	loan_amnt	term	int_rate	avg_cur_bal	in
count	1.832400e+04	1.832400e+04	17150.000000	18324.000000	18324.000000	18324.000000	18324.000000	17758.000000	9
mean	6.832645e+07	8.017611e+04	6.073178	467.543006	15522.661537	42.815979	13.850700	13466.600011	
std	4.245703e+07	6.487345e+04	3.639694	278.099801	9349.294243	10.822769	4.822253	16550.730832	
min	3.009180e+05	3.000000e+03	0.500000	30.650000	1000.000000	36.000000	5.310000	0.000000	
25%	3.491424e+07	4.700000e+04	2.000000	259.302500	8000.000000	36.000000	10.490000	3129.000000	
50%	6.838023e+07	6.500000e+04	6.000000	397.480000	14000.000000	36.000000	13.330000	7137.000000	
75%	9.730784e+07	9.500000e+04	10.000000	635.720000	21000.000000	60.000000	16.990000	18436.500000	
max	1.708249e+08	2.616000e+06	10.000000	1503.890000	40000.000000	60.000000	30.990000	341236.000000	

8 rows × 26 columns



```
data.isna()
```

```
#to see the number of missing values
```

	id	addr_state	annual_inc	emp_length	emp_title	home_ownership	installment	loan_amnt	purpose	term	...
0	False	False	False	False	False	False	False	False	False	False	...
1	False	False	False	False	False	False	False	False	False	False	...
2	False	False	False	False	False	False	False	False	False	False	...
3	False	False	False	False	False	False	False	False	False	False	...
4	False	False	False	False	False	False	False	False	False	False	...
...
18319	False	False	False	False	False	False	False	False	False	False	...
18320	False	False	False	False	False	False	False	False	False	False	...
18321	False	False	False	False	False	False	False	False	False	False	...
18322	False	False	False	False	False	False	False	False	False	False	...
18323	False	False	False	False	False	False	False	False	False	False	...

18324 rows × 31 columns

data.isna().sum()

```

id                0
addr_state        0
annual_inc        0
emp_length       1174
emp_title        1282
home_ownership    0
installment       0
loan_amnt         0
purpose           0
term              0
int_rate          0
avg_cur_bal       566
inq_last_12m     8929
max_bal_bc       8929
mo_sin_old_il_acct 1132
mo_sin_old_rev_tl_op 564

```

mo_sin_rcnt_rev_tl_op	564
mo_sin_rcnt_tl	564
mort_acc	398
mths_since_last_delinq	9048
num_bc_tl	564
num_il_tl	564
num_op_rev_tl	564
num_tl_90g_dpd_24m	564
num_tl_op_past_12m	564
open_acc	0
percent_bc_gt_75	610
pub_rec_bankruptcies	0
total_acc	0
total_bal_ex_mort	398
loan_status	0
dtype:	int64

```
set(data.duplicated())
```

```
{False}
```

```
# we can pass the result of the df.duplicated into a set to see if there is any instance of True
```

```
# No duplicates in the data
```

```
set(data.duplicated())
```

```
{False}
```

```
data.shape
```

```
(18324, 31)
```

```
# Checking the class imbalance
```

```
data['loan_status'].value_counts(normalize=True)
```


```
Fully Paid      0.786837
Charged Off     0.213163
Name: loan_status, dtype: float64
```

```
df=data
```

```
df.head()
```

	id	addr_state	annual_inc	emp_length	emp_title	home_ownership	installment	loan_amnt	purpose
0	802173	CA	72000.0	3.0	CA. Dept. Of Corrections	MORTGAGE	395.66	12000	debt_consolidation
1	14518910	TX	97500.0	1.0	Curriculum & Implementation Manager	RENT	966.47	35000	debt_consolidation
2	54333324	NY	120000.0	1.0	Senior manager	RENT	806.57	25000	credit_card
3	62247022	CA	130000.0	10.0	Border Patrol Agent	RENT	846.17	25225	debt_consolidation
4	71986114	TX	58296.0	10.0	Account Manager	MORTGAGE	41.79	1200	other

5 rows × 31 columns



```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18324 entries, 0 to 18323
Data columns (total 31 columns):
#   Column              Non-Null Count  Dtype
---  -
0   id                  18324 non-null  int64
```

1	addr_state	18324	non-null	object
2	annual_inc	18324	non-null	float64
3	emp_length	17150	non-null	float64
4	emp_title	17042	non-null	object
5	home_ownership	18324	non-null	object
6	installment	18324	non-null	float64
7	loan_amnt	18324	non-null	int64
8	purpose	18324	non-null	object
9	term	18324	non-null	int64
10	int_rate	18324	non-null	float64
11	avg_cur_bal	17758	non-null	float64
12	inq_last_12m	9395	non-null	float64
13	max_bal_bc	9395	non-null	float64
14	mo_sin_old_il_acct	17192	non-null	float64
15	mo_sin_old_rev_tl_op	17760	non-null	float64
16	mo_sin_rcnt_rev_tl_op	17760	non-null	float64
17	mo_sin_rcnt_tl	17760	non-null	float64
18	mort_acc	17926	non-null	float64
19	mths_since_last_delinq	9276	non-null	float64
20	num_bc_tl	17760	non-null	float64
21	num_il_tl	17760	non-null	float64
22	num_op_rev_tl	17760	non-null	float64
23	num_tl_90g_dpd_24m	17760	non-null	float64
24	num_tl_op_past_12m	17760	non-null	float64
25	open_acc	18324	non-null	int64
26	percent_bc_gt_75	17714	non-null	float64
27	pub_rec_bankruptcies	18324	non-null	int64
28	total_acc	18324	non-null	int64
29	total_bal_ex_mort	17926	non-null	float64
30	loan_status	18324	non-null	object

dtypes: float64(20), int64(6), object(5)


memory usage: 4.3+ MB

```
plt.figure(figsize=(15,10))
sns.histplot(df.loan_amnt)
plt.title('Loan Amount', fontsize = 25)
plt.show()
```



```
plt.figure(figsize=(12,8))
sns.histplot(df.loan_status)
plt.title('Distribution of Loan Status', fontsize = 25)
plt.show()
```


Distribution of Loan Status



df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18324 entries, 0 to 18323
Data columns (total 31 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     18324 non-null  int64
1   addr_state             18324 non-null  object
2   annual_inc             18324 non-null  float64
3   emp_length             17150 non-null  float64
4   emp_title              17042 non-null  object
```

5	home_ownership	18324	non-null	object
6	installment	18324	non-null	float64
7	loan_amnt	18324	non-null	int64
8	purpose	18324	non-null	object
9	term	18324	non-null	int64
10	int_rate	18324	non-null	float64
11	avg_cur_bal	17758	non-null	float64
12	inq_last_12m	9395	non-null	float64
13	max_bal_bc	9395	non-null	float64
14	mo_sin_old_il_acct	17192	non-null	float64
15	mo_sin_old_rev_tl_op	17760	non-null	float64
16	mo_sin_rcnt_rev_tl_op	17760	non-null	float64
17	mo_sin_rcnt_tl	17760	non-null	float64
18	mort_acc	17926	non-null	float64
19	mths_since_last_delinq	9276	non-null	float64
20	num_bc_tl	17760	non-null	float64
21	num_il_tl	17760	non-null	float64
22	num_op_rev_tl	17760	non-null	float64
23	num_tl_90g_dpd_24m	17760	non-null	float64
24	num_tl_op_past_12m	17760	non-null	float64
25	open_acc	18324	non-null	int64
26	percent_bc_gt_75	17714	non-null	float64
27	pub_rec_bankruptcies	18324	non-null	int64
28	total_acc	18324	non-null	int64
29	total_bal_ex_mort	17926	non-null	float64
30	loan_status	18324	non-null	object

dtypes: float64(20), int64(6), object(5)
memory usage: 4.3+ MB


Home ownership and loan status

df['home_ownership'].value_counts()

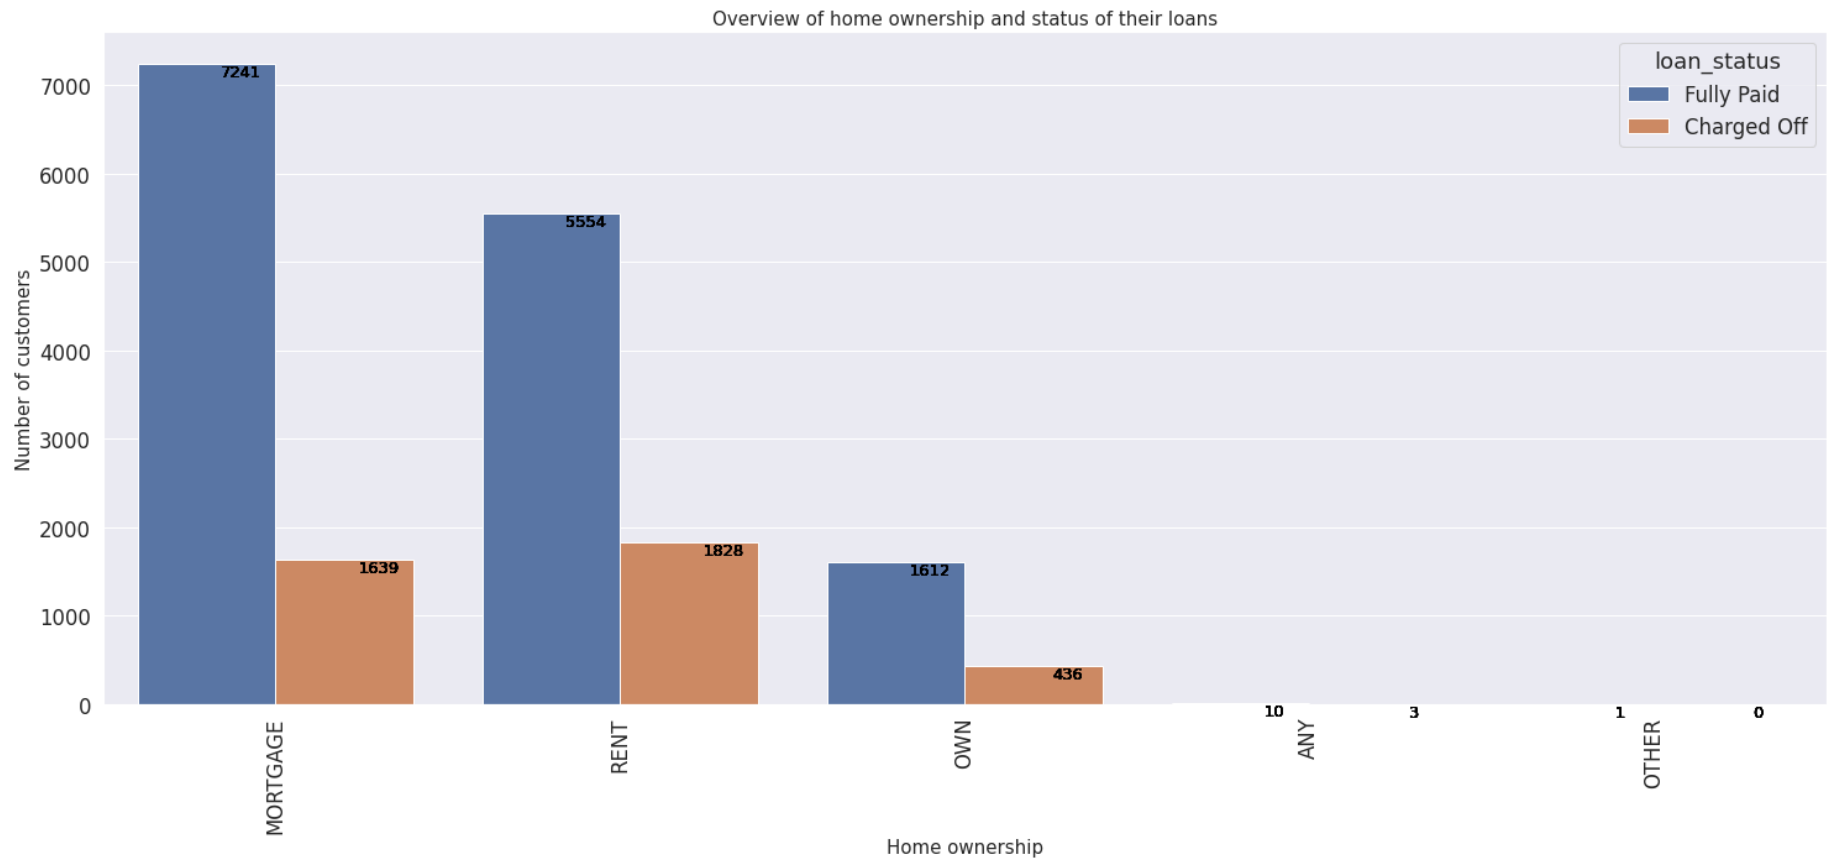
MORTGAGE	8880
RENT	7382
OWN	2048
ANY	13
OTHER	1

Name: home_ownership, dtype: int64

```
plt.figure(figsize=(20,10))
ax=sns.countplot(x='home_ownership', data=df)
ax.set_title('Distribution of home ownership' , fontsize = 15)
plt.xlabel('Home ownership', fontsize=15)
plt.ylabel('Number of customers', fontsize=15)
plt.xticks(rotation='vertical')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.5, p.get_height()), ha='center', va='top', color='white', size=13)
```




```
plt.figure(figsize=(25,10))
ax=sns.countplot(x='home_ownership',hue ='loan_status', data=df)
ax.set_title('Overview of home ownership and status of their loans' , fontsize = 15)
plt.xlabel('Home ownership', fontsize=15)
plt.ylabel('Number of customers', fontsize=15)
plt.xticks(rotation='vertical')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='top', color='black', size=13)
```




purpose and loan status

```
plt.figure(figsize=(20,10))
ax=sns.countplot(x='purpose', data=df)
ax.set_title('Purpose and loan status' , fontsize = 15)
plt.xlabel('Purpose', fontsize=15)
plt.xticks(rotation='vertical')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.5, p.get_height()), ha='center', va='top', color='black', size=13)
```

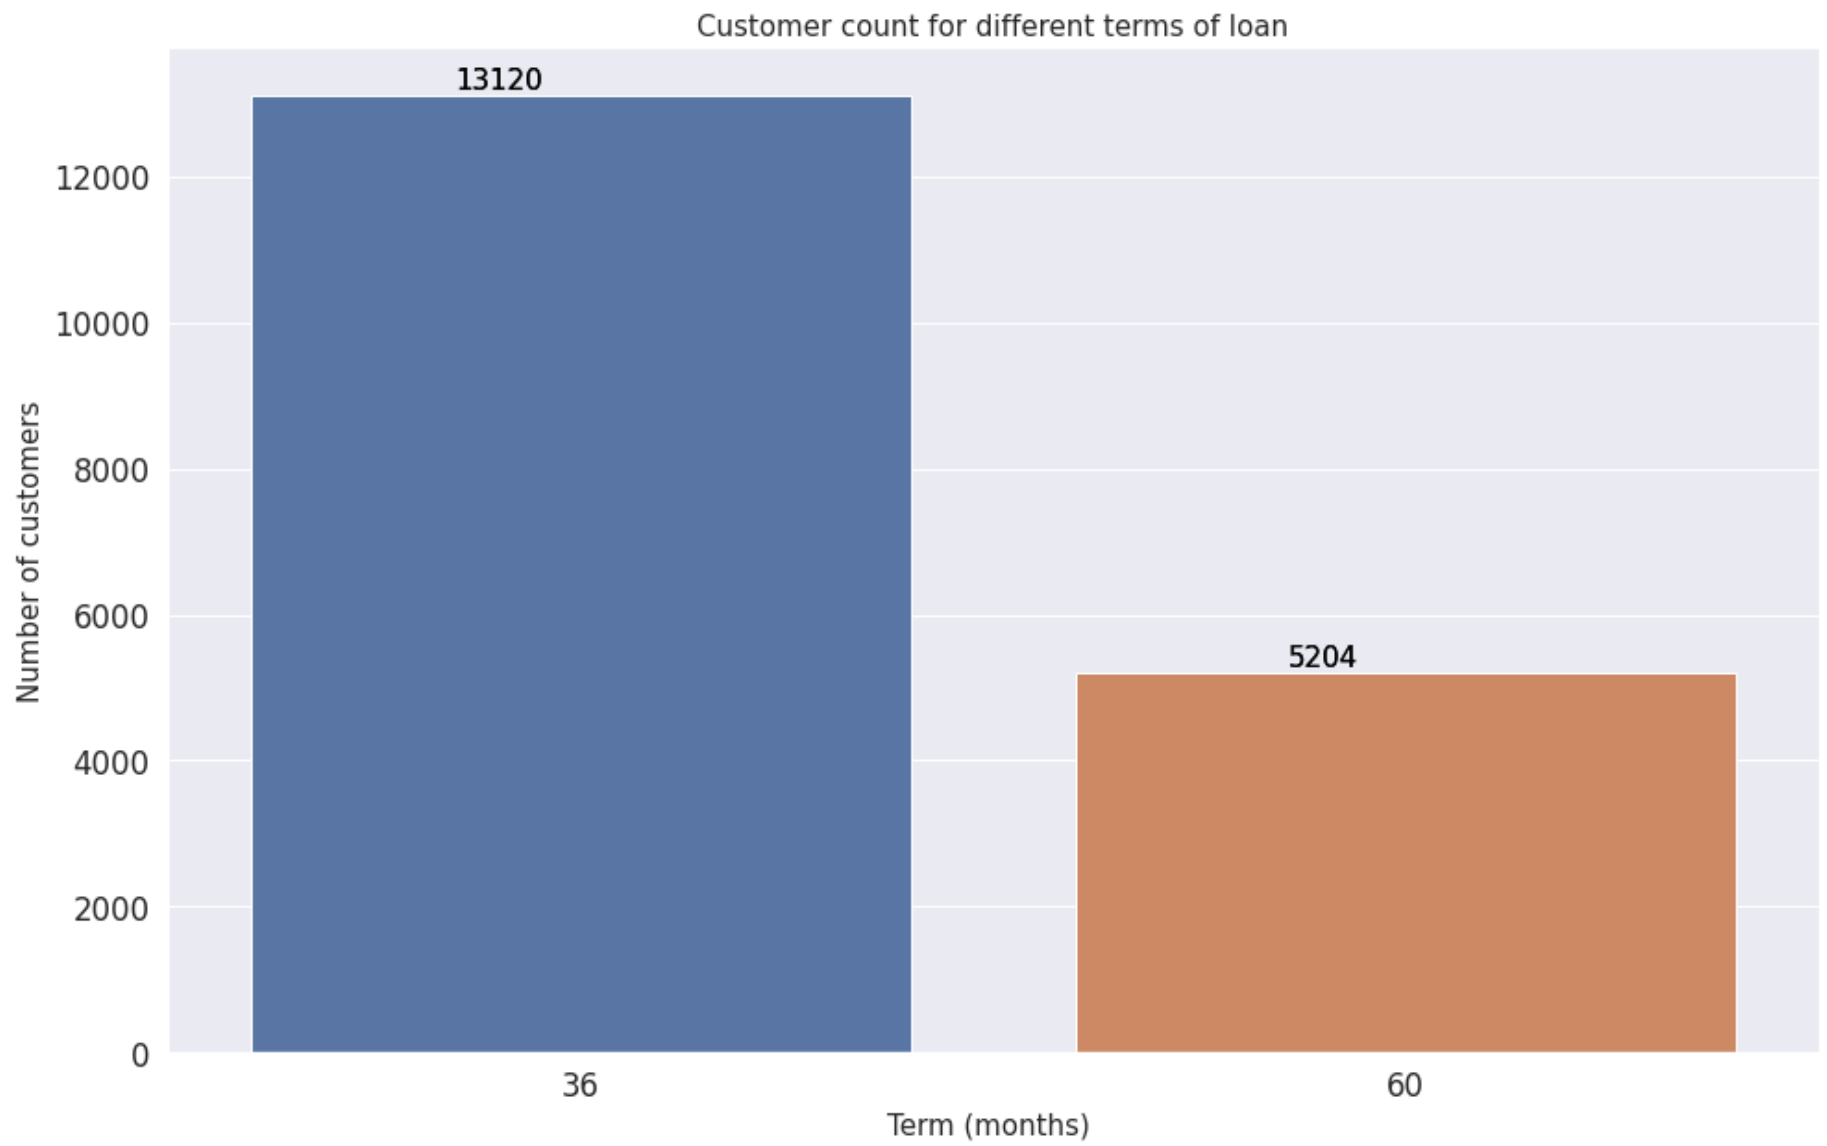


```
plt.figure(figsize=(28,14))
ax=sns.countplot(x='purpose', hue='loan_status', data=df)
ax.set_title('Purpose and loan status' , fontsize = 15)
plt.xlabel('Purpose', fontsize=15)
plt.xticks(rotation='vertical')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='center', color='black', size=13)
```



term and loan status

```
plt.figure(figsize=(16,10))
ax=sns.countplot(x='term', data=df)
ax.set_title('Customer count for different terms of loan' , fontsize = 15)
plt.xlabel('Term (months) ', fontsize=15)
plt.ylabel('Number of customers', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```

term and loan status

```
plt.figure(figsize=(20,11))
ax=sns.countplot(x='term', hue='loan_status', data=df)
ax.set_title('Term and loan status' , fontsize = 15)
plt.xlabel('Term (months)', fontsize=15)
plt.ylabel('Number of customers', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.25, p.get_height()), ha='center', va='top', color='black', size=15)
```

Term and loan status




```
# To put the loan amount in bins
bins = pd.cut(df['loan_amnt'], 5)
```

```
1m |
```

```
# loan status
```

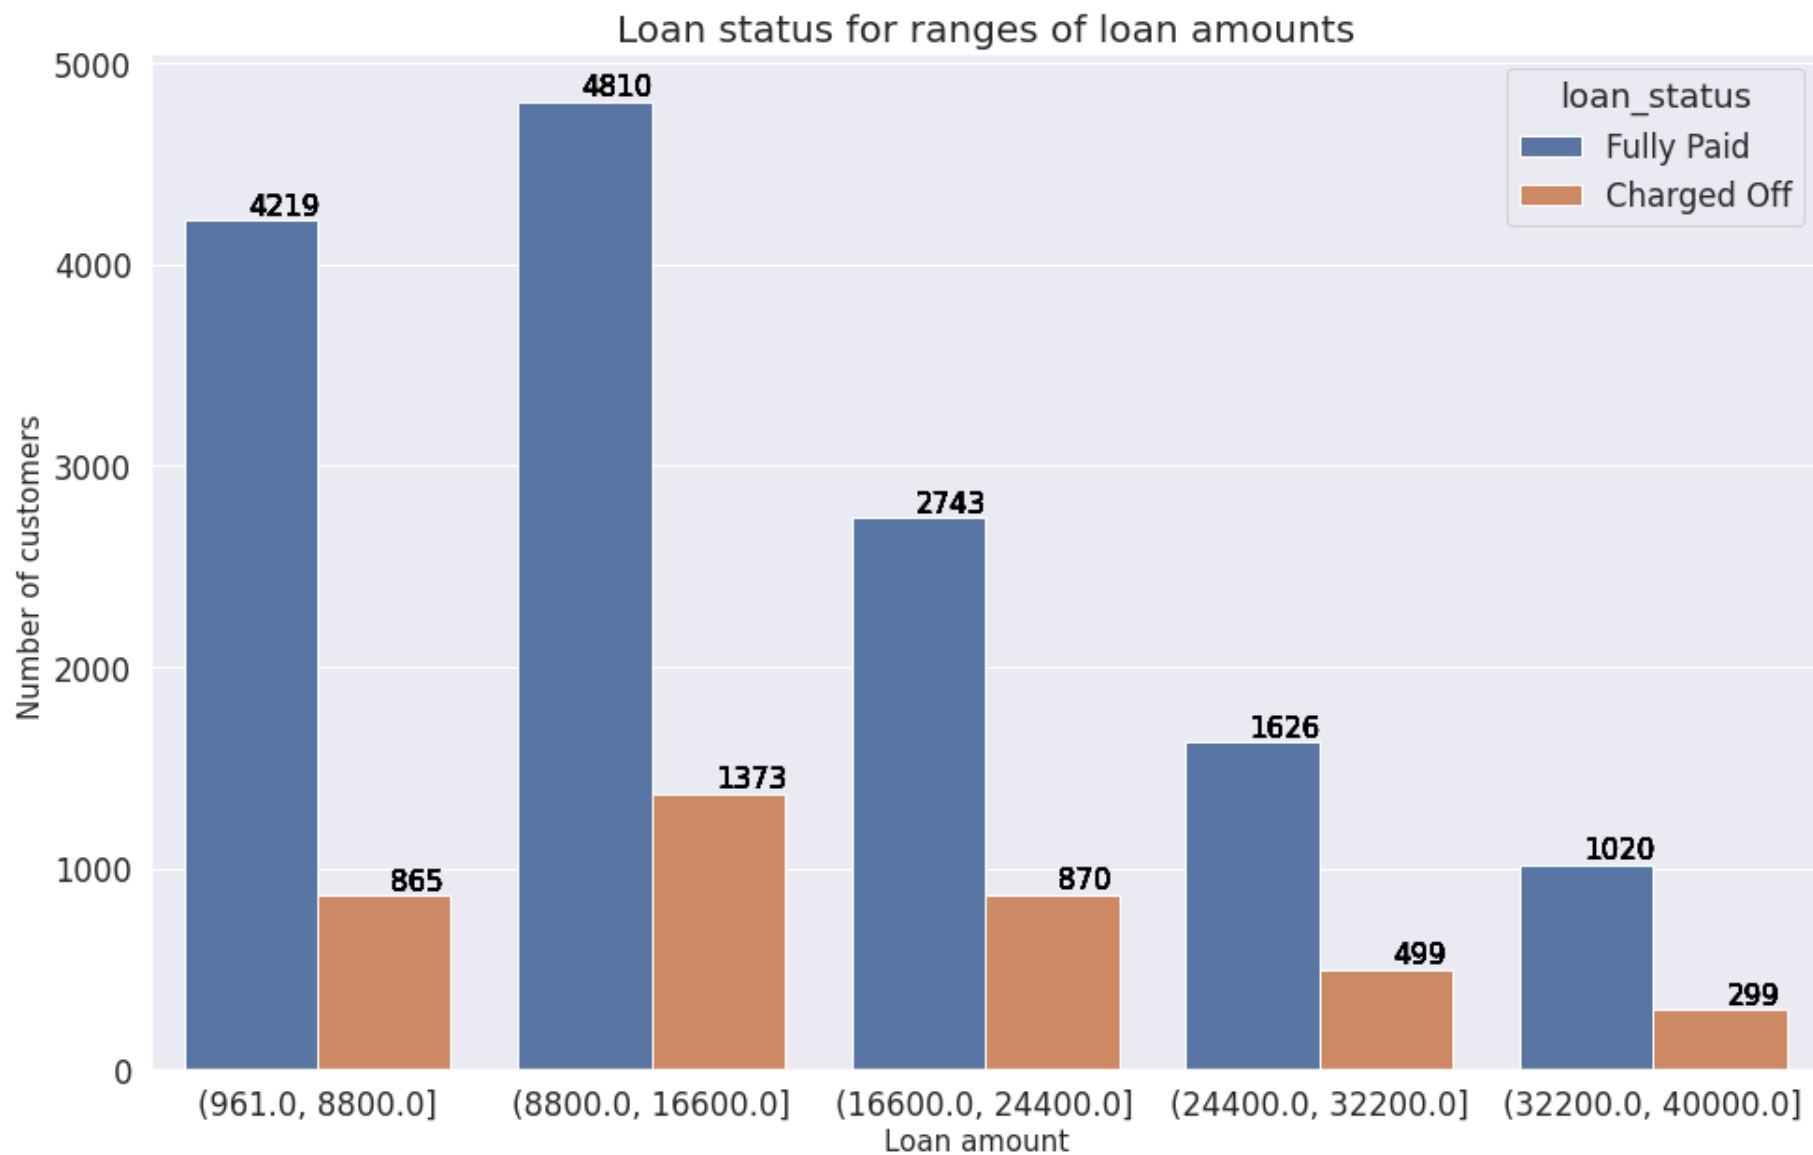
```
plt.figure(figsize=(20,10))
ax=sns.countplot(y=bins)
sns.set(font_scale=1.5)
ax.set_title('Count of customers depending on the amount of loan ', fontsize = 20)
plt.xlabel('Number of customers ', fontsize=15)
plt.ylabel('Range of loan amount ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```



```
# loan amount and loan status
```

```
plt.figure(figsize=(16,10))
sns.set(font_scale=1.5)
ax=sns.countplot(x=bins, hue='loan_status', data=df)
ax.set_title('Loan status for ranges of loan amounts' , fontsize = 20)
plt.xlabel('Loan amount ', fontsize=15)
plt.ylabel('Number of customers ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
```

```
ax.annotate(format(p.get_height(), '.0f'),  
            (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```




```

# Bin interest rate
# To put the loan amount in bins
bins2 = pd.cut(df['int_rate'], 6)

# interest rate


plt.figure(figsize=(20,10))
ax=sns.countplot(y=bins2)
sns.set(font_scale=1.5)
ax.set_title('Interest rate' , fontsize = 20)
plt.xlabel('Number of customers ', fontsize=15)
plt.ylabel('Range of interest rate ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)

```



interest rate and loan status

```
plt.figure(figsize=(20,10))
sns.set(font_scale=1.5)
ax=sns.countplot(x=bins2, hue='loan_status', data=df)
ax.set_title('Status of loans at different ranges of interest rates ', fontsize = 20)
plt.xlabel('Range of interest rate ', fontsize=15)
plt.ylabel('Number of customers ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```



```
# Bin average  
# To put the average current balance in bins  
bins3 = pd.cut(df['avg_cur_bal'], 10)
```

```
# average balance
```

```
plt.figure(figsize=(20,10))  
ax=sns.countplot(y=bins3)  
sns.set(font_scale=1.5)
```



```
ax.set_title('Average current balance' , fontsize = 20)
plt.xlabel('Number of customers ', fontsize=15)
plt.ylabel('Range of average balance ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```

Average current balance

(-341.236, 34123.6]

1

(68247.2, 102370.8]

Bin current balance

To put the average current balance in bins

bins4 = pd.cut(df['mths_since_last_delinq'], 10)

plt

average balance

plt.figure(figsize=(20,10))

ax=sns.countplot(y=bins4)

sns.set(font_scale=1.5)

ax.set_title('The number of months since last delinquency (missed payment)' , fontsize = 20)

plt.xlabel('Number of customers ', fontsize=15)

plt.ylabel('Range of months', fontsize=15)


plt.xticks(rotation='horizontal')

for p in ax.patches:


for p in ax.patches:

ax.annotate(format(p.get_height(), '.0f'),

(p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)



```
plt.figure(figsize=(20,10))
sns.set(font_scale=1.5)
ax=sns.countplot(x=bins4, hue='loan_status', data=df)
ax.set_title('The number of months since last delinquency (missed payment) ', fontsize = 20)
plt.xlabel('Range of months ', fontsize=15)
plt.ylabel('Number of customers ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```



```
df.nunique()
```

```
id          18324
addr_state    51
annual_inc   2434
emp_length    11
```

emp_title	10040
home_ownership	5
installment	10246
loan_amnt	1111
purpose	14
term	2
int_rate	465
avg_cur_bal	12635
inq_last_12m	25
max_bal_bc	6440
mo_sin_old_il_acct	355
mo_sin_old_rev_tl_op	564
mo_sin_rcnt_rev_tl_op	150
mo_sin_rcnt_tl	100
mort_acc	23
mths_since_last_delinq	98
num_bc_tl	43
num_il_tl	64
num_op_rev_tl	44
num_tl_90g_dpd_24m	11
num_tl_op_past_12m	20
open_acc	49
percent_bc_gt_75	111
pub_rec_bankruptcies	7
total_acc	93
total_bal_ex_mort	16450
loan_status	2
dtype: int64	

```
df.describe()
```

	id	annual_inc	emp_length	installment	loan_amnt	term	int_rate	avg_cur_bal	in
count	1.832400e+04	1.832400e+04	17150.000000	18324.000000	18324.000000	18324.000000	18324.000000	17758.000000	9
mean	6.832645e+07	8.017611e+04	6.073178	467.543006	15522.661537	42.815979	13.850700	13466.600011	
std	4.245703e+07	6.487345e+04	3.639694	278.099801	9349.294243	10.822769	4.822253	16550.730832	
min	3.009180e+05	3.000000e+03	0.500000	30.650000	1000.000000	36.000000	5.310000	0.000000	
50%	6.838023e+07	6.500000e+04	6.000000	397.480000	14000.000000	36.000000	13.330000	7137.000000	

df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 18324 entries, 0 to 18323
```


```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	id	18324 non-null	int64
1	addr_state	18324 non-null	object
2	annual_inc	18324 non-null	float64
3	emp_length	17150 non-null	float64
4	emp_title	17042 non-null	object
5	home_ownership	18324 non-null	object
6	installment	18324 non-null	float64
7	loan_amnt	18324 non-null	int64
8	purpose	18324 non-null	object
9	term	18324 non-null	int64
10	int_rate	18324 non-null	float64
11	avg_cur_bal	17758 non-null	float64
12	inq_last_12m	9395 non-null	float64
13	max_bal_bc	9395 non-null	float64
14	mo_sin_old_il_acct	17192 non-null	float64
15	mo_sin_old_rev_tl_op	17760 non-null	float64
16	mo_sin_rcnt_rev_tl_op	17760 non-null	float64
17	mo_sin_rcnt_tl	17760 non-null	float64
18	mort_acc	17926 non-null	float64
19	mths_since_last_delinq	9276 non-null	float64
20	num_bc_tl	17760 non-null	float64
21	num_il_tl	17760 non-null	float64

```
22  num_op_rev_tl          17760 non-null float64
23  num_tl_90g_dpd_24m     17760 non-null float64
24  num_tl_op_past_12m     17760 non-null float64
25  open_acc               18324 non-null int64
26  percent_bc_gt_75       17714 non-null float64
27  pub_rec_bankruptcies   18324 non-null int64
28  total_acc              18324 non-null int64
29  total_bal_ex_mort       17926 non-null float64
30  loan_status            18324 non-null object
dtypes: float64(20), int64(6), object(5)
memory usage: 4.3+ MB
```

```
plt.figure(figsize=(18,10))
df.boxplot(column=['annual_inc','emp_length','installment','term','avg_cur_bal','max_bal_bc'], return_type='axes')
```


<Axes: >



```
# now our data has no missing values, we can ahead to scale our data and train our model.  
# For this we will use KNN as the first model and Neural Networks for the second
```

```
plt.figure(figsize=(18,10))  
sns.boxplot(data=df[['annual_inc', 'avg_cur_bal', 'max_bal_bc']])
```


<Axes: >



#ML

data.describe()

```
# To drop the 'Id' column
data = data.drop('id', axis=1)
```

```
data.head()
```

	addr_state	annual_inc	emp_length	emp_title	home_ownership	installment	loan_amnt	purpose	term	int
0	CA	72000.0	3.0	CA. Dept. Of Corrections	MORTGAGE	395.66	12000	debt_consolidation	36	
1	TX	97500.0	1.0	Curriculum & Implementation Manager	RENT	966.47	35000	debt_consolidation	60	
2	NY	120000.0	1.0	Senior manager	RENT	806.57	25000	credit_card	36	
3	CA	130000.0	10.0	Border Patrol Agent	RENT	846.17	25225	debt_consolidation	36	
4	TX	58296.0	10.0	Account Manager	MORTGAGE	41.79	1200	other	36	

5 rows × 30 columns



```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18324 entries, 0 to 18323
Data columns (total 30 columns):
#   Column              Non-Null Count  Dtype
---  -
0   addr_state          18324 non-null  object
1   annual_inc          18324 non-null  float64
```

2	emp_length	17150	non-null	float64
3	emp_title	17042	non-null	object
4	home_ownership	18324	non-null	object
5	installment	18324	non-null	float64
6	loan_amnt	18324	non-null	int64
7	purpose	18324	non-null	object
8	term	18324	non-null	int64
9	int_rate	18324	non-null	float64
10	avg_cur_bal	17758	non-null	float64
11	inq_last_12m	9395	non-null	float64
12	max_bal_bc	9395	non-null	float64
13	mo_sin_old_il_acct	17192	non-null	float64
14	mo_sin_old_rev_tl_op	17760	non-null	float64
15	mo_sin_rcnt_rev_tl_op	17760	non-null	float64
16	mo_sin_rcnt_tl	17760	non-null	float64
17	mort_acc	17926	non-null	float64
18	mths_since_last_delinq	9276	non-null	float64
19	num_bc_tl	17760	non-null	float64
20	num_il_tl	17760	non-null	float64
21	num_op_rev_tl	17760	non-null	float64
22	num_tl_90g_dpd_24m	17760	non-null	float64
23	num_tl_op_past_12m	17760	non-null	float64
24	open_acc	18324	non-null	int64
25	percent_bc_gt_75	17714	non-null	float64
26	pub_rec_bankruptcies	18324	non-null	int64
27	total_acc	18324	non-null	int64
28	total_bal_ex_mort	17926	non-null	float64
29	loan_status	18324	non-null	object

dtypes: float64(20), int64(5), object(5)
memory usage: 4.2+ MB

data.isna().sum()

addr_state	0
annual_inc	0
emp_length	1174
emp_title	1282
home_ownership	0
installment	0
loan_amnt	0
purpose	0

term	0
int_rate	0
avg_cur_bal	566
inq_last_12m	8929
max_bal_bc	8929
mo_sin_old_il_acct	1132
mo_sin_old_rev_tl_op	564
mo_sin_rcnt_rev_tl_op	564
mo_sin_rcnt_tl	564
mort_acc	398
mths_since_last_delinq	9048
num_bc_tl	564
num_il_tl	564
num_op_rev_tl	564
num_tl_90g_dpd_24m	564
num_tl_op_past_12m	564
open_acc	0
percent_bc_gt_75	610
pub_rec_bankruptcies	0
total_acc	0
total_bal_ex_mort	398
loan_status	0
dtype:	int64

```
df2=data
```

```
# Drop missing values only in the 'emp_title' column
df2.dropna(subset=['emp_title'], inplace=True)
```

```
df2.isna().sum()
```

addr_state	0
annual_inc	0

emp_length	5
emp_title	0
home_ownership	0
installment	0
loan_amnt	0
purpose	0
term	0
int_rate	0
avg_cur_bal	528
inq_last_12m	8258
max_bal_bc	8258
mo_sin_old_il_acct	989
mo_sin_old_rev_tl_op	527
mo_sin_rcnt_rev_tl_op	527
mo_sin_rcnt_tl	527
mort_acc	369
mths_since_last_delinq	8335
num_bc_tl	527
num_il_tl	527
num_op_rev_tl	527
num_tl_90g_dpd_24m	527
num_tl_op_past_12m	527
open_acc	0
percent_bc_gt_75	552
pub_rec_bankruptcies	0
total_acc	0
total_bal_ex_mort	369
loan_status	0
dtype: int64	

```
df2.head()
```

	addr_state	annual_inc	emp_length	emp_title	home_ownership	installment	loan_amnt	purpose	term	int
0	CA	72000.0	3.0	CA. Dept. Of Corrections	MORTGAGE	395.66	12000	debt_consolidation	36	
1	TX	97500.0	1.0	Curriculum & Implementation Manager	RENT	966.47	35000	debt_consolidation	60	
2	NY	120000.0	1.0	Senior manager	RENT	806.57	25000	credit_card	36	
3	CA	130000.0	10.0	Border Patrol Agent	RENT	846.17	25225	debt_consolidation	36	

- - -



Identify columns with missing values

```
missing_cols = df2.columns[df2.isnull().any()]
```

missing_cols

```
Index(['emp_length', 'avg_cur_bal', 'inq_last_12m', 'max_bal_bc',
      'mo_sin_old_il_acct', 'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op',
      'mo_sin_rcnt_tl', 'mort_acc', 'mths_since_last_delinq', 'num_bc_tl',
      'num_il_tl', 'num_op_rev_tl', 'num_tl_90g_dpd_24m',
      'num_tl_op_past_12m', 'percent_bc_gt_75', 'total_bal_ex_mort'],
      dtype='object')
```

```
li=list(missing_cols)
```

li

```
['emp_length',
 'avg_cur_bal',
```

```
'inq_last_12m',
'max_bal_bc',
'mo_sin_old_il_acct',
'mo_sin_old_rev_tl_op',
'mo_sin_rcnt_rev_tl_op',
'mo_sin_rcnt_tl',
'mort_acc',
'mths_since_last_delinq',
'num_bc_tl',
'num_il_tl',
'num_op_rev_tl',
'num_tl_90g_dpd_24m',
'num_tl_op_past_12m',
'percent_bc_gt_75',
'total_bal_ex_mort']
```

```
# replace missing values using KNN imputation
# instantitating KNN imputer
imputer = KNNImputer(n_neighbors=5)
```

```
df2[li] = imputer.fit_transform(df2[li])
```

```
df2.isna().sum()
```

addr_state	0
annual_inc	0
emp_length	0
emp_title	0
home_ownership	0
installment	0
loan_amnt	0
purpose	0
term	0
int_rate	0
avg_cur_bal	0
inq_last_12m	0

```
max_bal_bc      0
mo_sin_old_il_acct  0
mo_sin_old_rev_tl_op  0
mo_sin_rcnt_rev_tl_op  0
mo_sin_rcnt_tl    0
mort_acc        0
mths_since_last_delinq  0
num_bc_tl       0
num_il_tl       0
num_op_rev_tl   0
num_tl_90g_dpd_24m  0
num_tl_op_past_12m  0
open_acc        0
percent_bc_gt_75  0
pub_rec_bankruptcies  0
total_acc       0
total_bal_ex_mort  0
loan_status     0
dtype: int64
```

```
df3=df2
```

```
df3.shape
```

```
(17042, 30)
```

```
df3['loan_status'].value_counts(normalize=True)
```

```
Fully Paid      0.789872
Charged Off     0.210128
Name: loan_status, dtype: float64
```

```
# handle class imbalance using oversampling
X = df3.drop('loan_status', axis=1)
y = df3['loan_status']
```



```
sampler = RandomOverSampler()  
X_resampled, y_resampled = sampler.fit_resample(X, y)
```

```
X_resampled.shape
```

```
(26922, 29)
```

```
y_resampled.shape
```


```
(26922,)
```

```
y_resampled.value_counts(normalize=True)
```

```
Fully Paid      0.5  
Charged Off     0.5  
Name: loan_status, dtype: float64
```

```
plt.figure(figsize=(15,8))  
plt.title('Class Balance', fontsize=15)  
sns.countplot(x= y_resampled)
```

```
<Axes: title={'center': 'Class Balance'}, xlabel='loan_status', ylabel='count'>
```



```
#Encoding
```

```
# encode categorical variables
```

```
le = LabelEncoder()
```

```
for col in X_resampled.select_dtypes(include='object'):
```

```
    X_resampled[col] = le.fit_transform(X_resampled[col])
```



```
X_resampled.head()
```

	addr_state	annual_inc	emp_length	emp_title	home_ownership	installment	loan_amnt	purpose	term	int_rate	...
0	4	72000.0	3.0	1063	1	395.66	12000	2	36	11.49	...
1	43	97500.0	1.0	2003	4	966.47	35000	2	60	21.00	...

```
y_resampled.tail()
```

```

26917    Charged Off
26918    Charged Off
26919    Charged Off
26920    Charged Off
26921    Charged Off
Name: loan_status, dtype: object

```



```

# Replace categorical values in the outcome column
y_resampled = y_resampled.replace({'Fully Paid': 1, 'Charged Off': 0})

```

```

# Check the new values in the outcome column
print(y_resampled.unique())

```

```
[1 0]
```

```
y_resampled.tail()
```

```

26917    0
26918    0
26919    0
26920    0
26921    0
Name: loan_status, dtype: int64

```

```
X = X_resampled
```

```
y = y_resampled
```

```
corr = df3.corr()  
corr
```

	annual_inc	emp_length	installment	loan_amnt	term	int_rate	avg_cur_bal	inq_last_12m	m
annual_inc	1.000000	0.068835	0.377717	0.384595	0.038419	-0.104523	0.352539	0.080676	
emp_length	0.068835	1.000000	0.080613	0.096398	0.062969	-0.002952	0.093442	0.005957	
installment	0.377717	0.080613	1.000000	0.948106	0.120482	0.120168	0.209443	0.045250	
loan_amnt	0.384595	0.096398	0.948106	1.000000	0.372737	0.104367	0.232020	0.039641	
term	0.038419	0.062969	0.120482	0.372737	1.000000	0.401595	0.055051	0.041989	
int_rate	-0.104523	-0.002952	0.120168	0.104367	0.401595	1.000000	-0.088488	0.119741	
avg_cur_bal	0.352539	0.093442	0.209443	0.232020	0.055051	-0.088488	1.000000	0.059283	

```
#to get variables that are highly correlated
plt.figure(figsize=(25,15))

# Calculate the correlation matrix
corr = df3.corr()

# Create a heatmap to visualize the correlation matrix
sns.heatmap(corr, cmap='coolwarm', annot=True, fmt='.2f', square=True)


# Set a threshold for correlation coefficient
threshold = 0.5

# Find the highly correlated features
highly_correlated = []
for i in range(len(corr.columns)):
    for j in range(i):
        if abs(corr.iloc[i, j]) > threshold:
            colname = corr.columns[i]
            highly_correlated.append(colname)

# Drop the highly correlated features
#df.drop(highly_correlated, axis=1, inplace=True)
print(highly_correlated)
```



```
['loan_amnt', 'mo_sin_rcnt_tl', 'num_op_rev_tl', 'open_acc', 'open_acc', 'total_acc', 'total_acc', 'total_acc', 'total_
```



```
print(highly_correlated)
```

```
['loan_amnt', 'mo_sin_rcnt_tl', 'num_op_rev_tl', 'open_acc', 'open_acc', 'total_acc', 'total_acc', 'total_acc', 'total_
```



```
# Instantiate SelectKBest with f_classif as the scoring function
```

```
selector = SelectKBest(score_func=f_classif, k=10)
```

```
# Fit the selector to the data
```

```
selector.fit(X_resampled, y_resampled)
```

```
# Get the indices of the selected features
```

```
selected_features_indices = selector.get_support(indices=True)
```

```

# Get the names of the selected features
selected_features_names = X_resampled.columns[selected_features_indices]

# Print the names of the selected features
print(selected_features_names)

Index(['annual_inc', 'home_ownership', 'loan_amnt', 'term', 'int_rate',
      'avg_cur_bal', 'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op',
      'mort_acc', 'num_tl_op_past_12m'],
      dtype='object')

# Display the scores of the top 5 features
scores = selector.scores_
top_k_scores = sorted(scores, reverse=True)[:10]
top_k_indices = np.argsort(scores)[::-1][:10]

print("Top 5 feature scores:")
for i in range(len(top_k_scores)):
    print("Feature {}: Score = {:.2f}".format(top_k_indices[i], top_k_scores[i]))

```

```

Top 5 feature scores:
Feature 9: Score = 2463.85
Feature 8: Score = 1089.43
Feature 10: Score = 342.67
Feature 17: Score = 324.50
Feature 23: Score = 247.17
Feature 4: Score = 223.77
Feature 14: Score = 138.89
Feature 15: Score = 132.45
Feature 6: Score = 126.70
Feature 1: Score = 95.74

```

```

feature_names = df3.drop('loan_status', axis=1).columns

```




```
# Get the names and scores of the top 10 features
feature_names = df3.drop('loan_status', axis=1).columns
top_scores = selector.scores_.argsort()[-10:][::-1]
top_features = feature_names[top_scores]

# Print the names and scores of the top 10 features
for i, feature in enumerate(top_features):
    print("{} . {} ({:.2f})".format(i+1, feature, selector.scores_[top_scores][i]))

# Create a bar plot of the top 10 features and their scores
plt.figure(figsize=(20,8))
sns.set(font_scale=1.5)
plt.bar(range(len(top_scores)), selector.scores_[top_scores])
plt.xticks(range(len(top_scores)), top_features, rotation='vertical')
plt.xlabel("Feature")
plt.ylabel("Score")
plt.title("Top 10 Features selected by SelectKBest")
plt.show()
```

1. int_rate (2463.85)
2. term (1089.43)
3. avg_cur_bal (342.67)
4. mort_acc (324.50)
5. num_tl_op_past_12m (247.17)
6. home_ownership (223.77)
7. mo_sin_old_rev_tl_op (138.89)
8. mo_sin_rcnt_rev_tl_op (132.45)
9. loan_amnt (126.70)
10. annual_inc (95.74)



```
df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17042 entries, 0 to 18323
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---

```

0	addr_state	17042	non-null	object
1	annual_inc	17042	non-null	float64
2	emp_length	17042	non-null	float64
3	emp_title	17042	non-null	object
4	home_ownership	17042	non-null	object
5	installment	17042	non-null	float64
6	loan_amnt	17042	non-null	int64
7	purpose	17042	non-null	object
8	term	17042	non-null	int64
9	int_rate	17042	non-null	float64
10	avg_cur_bal	17042	non-null	float64
11	inq_last_12m	17042	non-null	float64
12	max_bal_bc	17042	non-null	float64
13	mo_sin_old_il_acct	17042	non-null	float64
14	mo_sin_old_rev_tl_op	17042	non-null	float64
15	mo_sin_rcnt_rev_tl_op	17042	non-null	float64
16	mo_sin_rcnt_tl	17042	non-null	float64
17	mort_acc	17042	non-null	float64
18	mths_since_last_delinq	17042	non-null	float64
19	num_bc_tl	17042	non-null	float64
20	num_il_tl	17042	non-null	float64
21	num_op_rev_tl	17042	non-null	float64
22	num_tl_90g_dpd_24m	17042	non-null	float64
23	num_tl_op_past_12m	17042	non-null	float64
24	open_acc	17042	non-null	int64
25	percent_bc_gt_75	17042	non-null	float64
26	pub_rec_bankruptcies	17042	non-null	int64
27	total_acc	17042	non-null	int64
28	total_bal_ex_mort	17042	non-null	float64
29	loan_status	17042	non-null	object

dtypes: float64(20), int64(5), object(5)

memory usage: 4.0+ MB

```
# select the top K features using f_classic
kbest = SelectKBest(score_func=f_classif, k=10)
X_resampled = kbest.fit_transform(X_resampled, y_resampled)

# split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.3, random_state = 40)
```

```
scaler = StandardScaler()  
X_train_sc = scaler.fit_transform(X_train)  
X_test_sc = scaler.transform(X_test)
```

```
# train a decision tree classifier on the data  
clf = DecisionTreeClassifier()
```

```
#clf = RandomForestClassifier()
```

```
clf.fit(X_train_sc, y_train)
```

```
# test the classifier on the test set and print the classification report  
y_pred = clf.predict(X_test_sc)  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.79	0.95	0.86	4053
1	0.94	0.75	0.83	4024
accuracy			0.85	8077
macro avg	0.86	0.85	0.85	8077
weighted avg	0.86	0.85	0.85	8077

```
Accuracy = metrics.accuracy_score(y_test, y_pred)  
print('Accuracy score: %.2f\n\n'%(Accuracy))
```

```

conf_matrix = metrics.confusion_matrix(y_test, y_pred)
print('The confusion matrix is:')
print(conf_matrix, '\n\n')
print('-----')
result = metrics.classification_report(y_test, y_pred)
print('Classification Report:\n')
print(result)

```

Accuracy score:0.85

The confusion matrix is:

```

[[3844  209]
 [1001 3023]]

```


Classification Report:

	precision	recall	f1-score	support
0	0.79	0.95	0.86	4053
1	0.94	0.75	0.83	4024
accuracy			0.85	8077
macro avg	0.86	0.85	0.85	8077
weighted avg	0.86	0.85	0.85	8077

```

plt.figure(figsize=(15,8))
zx = sns.heatmap(conf_matrix, cmap='flare', annot_kws={"size": 15}, annot= True, fmt = 'd')
plt.title('Confusion Matrix using Decision Tree Classifier ', fontsize= 15)
plt.xlabel('Predicted Ouput', fontsize =15)
plt.ylabel('Actual Output', fontsize =15)
plt.show()

```



```
#Tuning decision tree
#from sklearn.tree import DecisionTreeClassifier
#from sklearn.model_selection import GridSearchCV

# Define the Decision Tree classifier
dt = DecisionTreeClassifier()

# Define the hyperparameters to tune
param_grid = {'max_depth': [2, 4, 6, 8, 10],
              'min_samples_split': [2, 4, 6, 8, 10],
              'min_samples_leaf': [1, 2, 3, 4, 5]}

# Perform hyperparameter tuning using GridSearchCV
```

```
grid_search = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, verbose=0)
grid_search.fit(X, y)
```

```
# Print the results
print("Best accuracy score: {:.2f}".format(grid_search.best_score_))
print("Best parameters: {}".format(grid_search.best_params_))
```

```
Best accuracy score: 0.70
Best parameters: {'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 4}
```

```
#Using RandomForest
```

```
# train a decision tree classifier on the data
#clf = DecisionTreeClassifier()
```

```
clf = RandomForestClassifier()
```

```
clf.fit(X_train_sc, y_train)
```

```
# test the classifier on the test set and print the classification report
y_pred = clf.predict(X_test_sc)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.95	0.92	4053
1	0.94	0.88	0.91	4024
accuracy			0.91	8077
macro avg	0.92	0.91	0.91	8077

weighted avg	0.92	0.91	0.91	8077
--------------	------	------	------	------

```
Accuracy = metrics.accuracy_score(y_test, y_pred)
print('Accuracy score:%.2f\n\n'%(Accuracy))
conf_matrix = metrics.confusion_matrix(y_test, y_pred)
print('The confusion matrix is:')
print(conf_matrix, '\n\n')
print('-----')
result = metrics.classification_report(y_test, y_pred)
print('Classification Report:\n')
print(result)
```

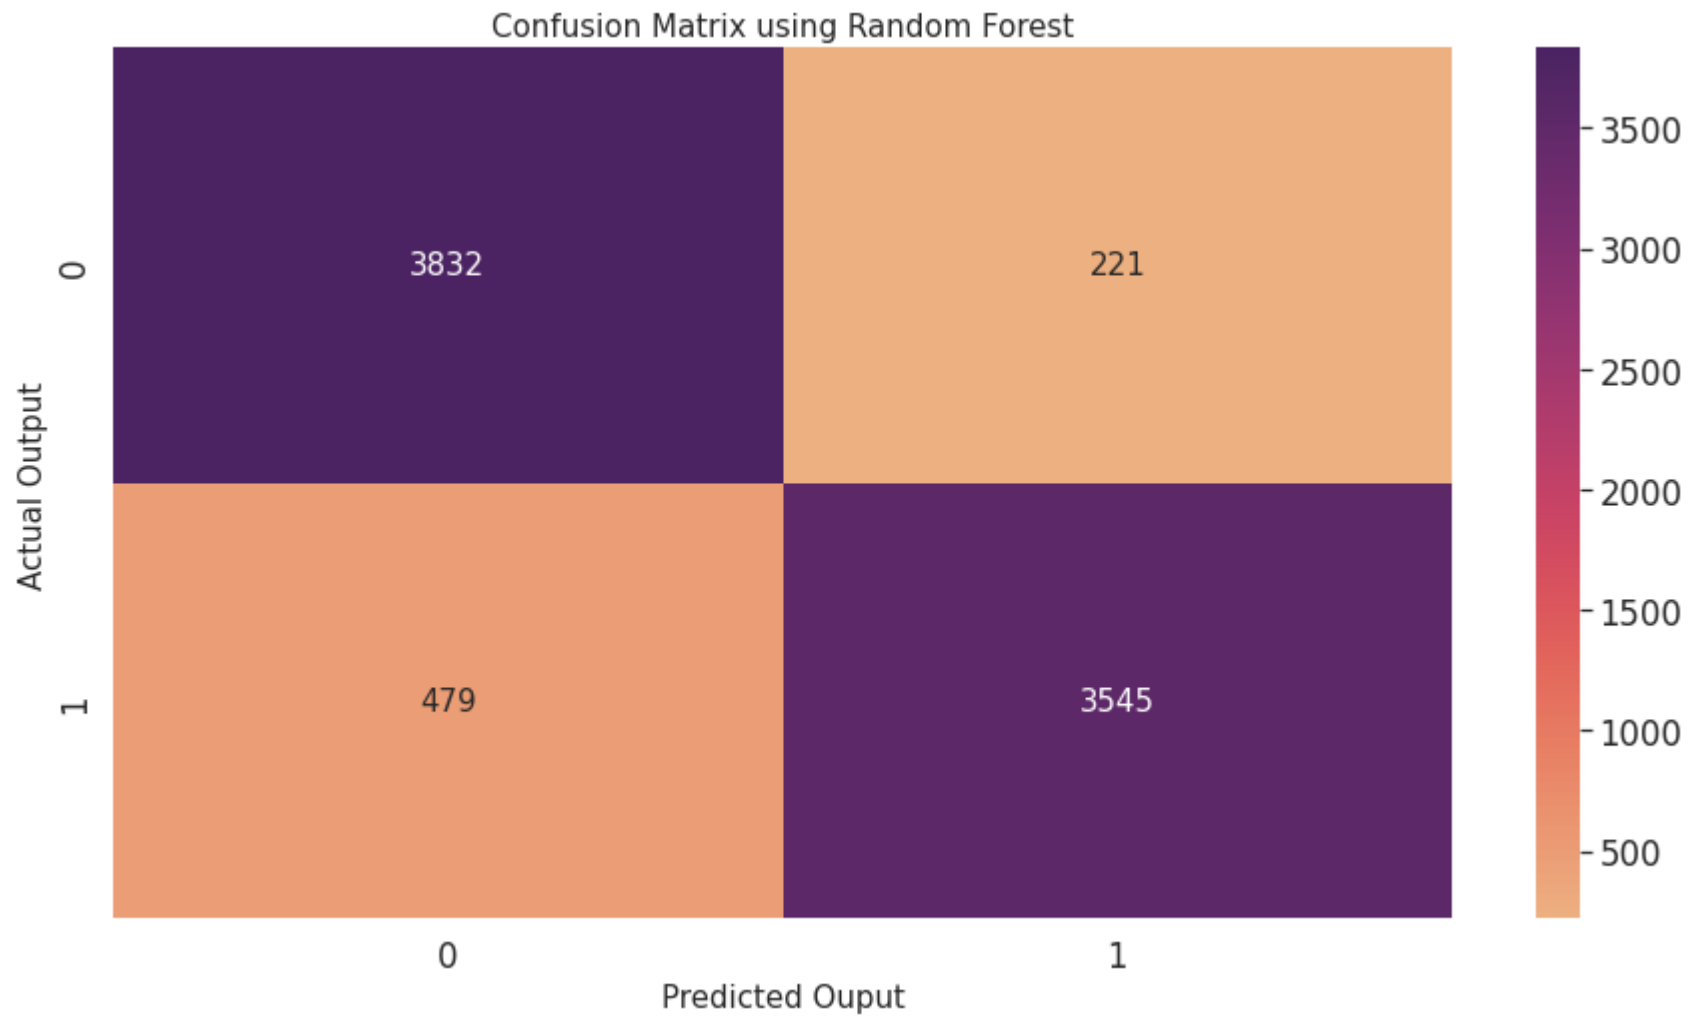
Accuracy score:0.91

The confusion matrix is:
[[3832 221]
[479 3545]]

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.95	0.92	4053
1	0.94	0.88	0.91	4024
accuracy			0.91	8077
macro avg	0.92	0.91	0.91	8077
weighted avg	0.92	0.91	0.91	8077


```
plt.figure(figsize=(15,8))
zx = sns. heatmap(conf_matrix, cmap = 'flare',annot_kws={"size": 15}, annot= True, fmt = 'd')
plt.title('Confusion Matrix using Random Forest', fontsize= 15)
plt.xlabel('Predicted Ouput', fontsize =15)
plt.ylabel('Actual Output', fontsize =15)
plt.show()
```



```
# GRADIENT BOOSTED DECISION TREE
```

```
#Using Gradient Boosted Decision Tree
```

```
# train a decision tree classifier on the data
```

```
#clf = DecisionTreeClassifier()
```

```
#clf = RandomForestClassifier()
```

```
clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.7, max_depth=16, random_state=42)
```

```
clf.fit(X_train_sc, y_train)
```

```
# test the classifier on the test set and print the classification report
```

```
y_pred = clf.predict(X_test_sc)
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.92	0.94	0.93	4053
1	0.93	0.92	0.93	4024
accuracy			0.93	8077
macro avg	0.93	0.93	0.93	8077
weighted avg	0.93	0.93	0.93	8077

```
from sklearn import metrics
```

```
Accuracy = metrics.accuracy_score(y_test, y_pred)
```

```
print('Accuracy score: %.2f\n\n'%(Accuracy))
```

```
conf_matrix = metrics.confusion_matrix(y_test, y_pred)
```

```
print('The confusion matrix is:')
```

```
print(conf_matrix, '\n\n')
```

```
print('-----')
```

```
result = metrics.classification_report(y_test, y_pred)
```

```
print('Classification Report:\n')
```

```
print(result)
```

Accuracy score:0.93


The confusion matrix is:

```
[[3790 263]
 [ 327 3697]]
```

```
-----
Classification Report:
```

	precision	recall	f1-score	support
0	0.92	0.94	0.93	4053
1	0.93	0.92	0.93	4024
accuracy			0.93	8077
macro avg	0.93	0.93	0.93	8077
weighted avg	0.93	0.93	0.93	8077

```
plt.figure(figsize=(20,8))
zx = sns. heatmap(conf_matrix, cmap = 'flare', annot_kws={"size": 18},annot= True, fmt = 'd')
plt.title('Confusion Matrix using Gradient Boosted Decision Tree ', fontsize= 15)
plt.xlabel('Predicted Ouput', fontsize =15)
plt.ylabel('Actual Output', fontsize =15)
plt.show()
```



F1-score of 94%

Identifying top 10 features driving the Gradient Boosted Decision model by allowing the model to select importance feature:

split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state = 40)
```

```
scaler = StandardScaler()
```

```
X_train_scale = scaler.fit_transform(X_train)
```

```
X_test_scale = scaler.transform(X_test)
```

Step 4: Model Evaluation

```
from sklearn.metrics import f1_score
```

```

scaler = StandardScaler()
X_train_scale = scaler.fit_transform(X_train)
X_test_scale = scaler.transform(X_test)

clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.7, max_depth=16, random_state=42)
clf.fit(X_train_scale, y_train)
y_pred = clf.predict(X_test_scale)
score = f1_score(y_test, y_pred)
print(f"F1-score: {score:.4f}")

```

F1-score: 0.9350

```

from sklearn import metrics
Accuracy = metrics.accuracy_score(y_test, y_pred)
print('Accuracy score: %.2f\n\n' % (Accuracy))
conf_matrix = metrics.confusion_matrix(y_test, y_pred)
print('The confusion matrix is:')
print(conf_matrix, '\n\n')
print('-----')
results = metrics.classification_report(y_test, y_pred)
print('Classification Report:\n')
print(results)

```

Accuracy score: 0.94

The confusion matrix is:

```

[[3787  266]
 [ 258 3766]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.93	0.94	4053
1	0.93	0.94	0.93	4024

accuracy			0.94	8077
macro avg	0.94	0.94	0.94	8077
weighted avg	0.94	0.94	0.94	8077

```
plt.figure(figsize=(20,8))
zx = sns. heatmap(conf_matrix, cmap = 'flare', annot_kws={"size": 18},annot= True, fmt = 'd')
plt.title('Confusion Matrix using Gradient Boosted Decision Tree ', fontsize= 15)
plt.xlabel('Predicted Ouput', fontsize =15)
plt.ylabel('Actual Output', fontsize =15)
plt.show()
```

Confusion Matrix using Gradient Boosted Decision Tree




#Step 5: Feature Importance Analysis

```
feature_importances = clf.feature_importances_  
feature_names = df3.drop('loan_status', axis=1).columns
```



Step 6: Plot Feature Importance Graph

```
top_features = pd.Series(feature_importances, index=feature_names).sort_values(ascending=False)[:5]  
plt.figure(figsize=(27,8))  
plt.bar(top_features.index, top_features)  
plt.title('Top 5 Features Driving Model')  
plt.xlabel('Feature Name')  
plt.ylabel('Feature Importance')  
plt.show()
```



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 10:34PM

